

Creative or Not? Hierarchical Diffusion Modeling of the Creative Evaluation Process

Michelle C. Donzallaz¹, Julia M. Haaf¹, & Claire E. Stevenson¹

¹ University of Amsterdam

Draft version 1, April 2021. This paper has not been peer reviewed. Please do not copy or cite without authors' permission.

Author Note

This report was written in R-Markdown with code for data analysis integrated into the text. The source code and the data can be found at: <https://osf.io/73c2d/>. The research reported in this manuscript has been presented at the virtual MathPsych/ICCM conference in 2020.

Correspondence concerning this article should be addressed to Claire E. Stevenson, Nieuwe Achtergracht 129-B, 1018 WS Amsterdam. E-mail: c.e.stevenson@uva.nl

Abstract

When producing creative ideas (i.e., ideas that are original and useful) two main processes occur: ideation, where people brainstorm ideas, and evaluation, where they decide if the ideas are creative or not. While much is known about the ideation phase, the cognitive processes involved in creativity evaluation are largely unclear. In this paper, we present a novel modeling approach for the evaluation phase of creativity. We apply the drift diffusion model (DDM) to the creative-or-not (CON)-task to study the cognitive basis of evaluation and to examine individual differences in the extent to which people take originality and utility into account when evaluating creative ideas. The CON-task is a timed decision-making task where participants indicate whether they find uses for certain objects creative or not (e.g., using a book as a buoy). The different use items vary on the two creativity dimensions ‘originality’ and ‘utility’. In two studies ($n = 293$, 17806 trials; $n = 152$, 9291 trials), we found that stimulus originality was strongly related to participants’ drift rate, whereas stimulus utility was only somewhat associated with the drift rate. However, participants differed substantially in the effects of originality and utility. Furthermore, the implicit weights assigned to originality and utility on the CON-task were aligned with self-reported importance ratings of originality and utility and associated with divergent thinking performance. Our findings underline the importance of communicating rating criteria in divergent thinking tasks such as the alternative uses task to ensure a fair assessment of creative ability.

Keywords: creativity, evaluation, diffusion model, Bayesian hierarchical modeling

Creative or Not? Hierarchical Diffusion Modeling of the Creative Evaluation Process

Creative ideas are essential for tackling today's problems, from personal obstacles, such as combining work and childcare during a pandemic lockdown, to societal challenges such as climate change (Hennessey & Amabile, 2010). It is no wonder that educators are increasingly including creativity in curricula and that managers consider creativity a key skill (Casner-Lotto & Barrington, 2006; IBM, 2010). Divergent thinking tasks are often used to assess creative ability, the ability to produce original and useful ideas (Barron, 1955; Runco & Jaeger, 2012; Stein, 1953). Perhaps the most common divergent thinking measure is the alternative uses task (AUT; Guilford, 1967; Runco & Acar, 2012). On the AUT, people are typically asked to come up with as many unusual uses as possible for a given object (e.g., "bath toy" for the object "brick") within a certain time interval. When solving the AUT, two main processes occur: ideation and evaluation (Basadur, 1995; Guilford, 1967; Runco & Acar, 2012). Ideation is the "brainstorm" phase where one comes up with ideas, and evaluation is the decision making phase where one judges which ideas are creative enough to pursue (or in the case of the AUT to list as a response). While most research has focused on ideation and how to improve it (e.g., Benedek, Fink, & Neubauer, 2006; Forthmann et al., 2019), how people evaluate and select ideas is less understood (de Buisonjé, Ritter, de Bruin, ter Horst, & Meeldijk, 2017; Herman & Reiter-Palmon, 2011; Hunter, Friedrich, Bedell, & Mumford, 2006; Rietzschel, Nijstad, & Stroebe, 2010; Runco & Smith, 1992). Yet, in the real world, it only takes one well-selected idea to solve a problem (e.g., the printing press reduced the cost and labor of printing books, thereby increasing literacy and making knowledge more accessible) or disregarding a good idea to lose the battle in innovative business (e.g., Blockbuster vs. Netflix; Randolph, 2019). In this paper, we take a cognitive modeling approach to fill this gap and study how people decide which ideas are creative or not on the AUT.

While the standard definition of creativity states that originality and utility are both

needed for creativity, individuals may differ in how much they value these two dimensions. Previous research has mainly focused on population-level effects and found that people tend to value originality more than utility when judging the creativity of an idea (Caroff & Besançon, 2008; Diedrich, Benedek, Jauk, & Neubauer, 2015; Runco & Charles, 1993). However, creativity is a much-used word and frequently discussed by laypeople (Mueller, Melwani, Loewenstein, & Deal, 2018). Consequently, different people may have substantially different conceptions of the importance of originality and utility for creativity (e.g., Loewenstein & Mueller, 2016). For example, some may value utility more than others when judging creativity, and some may even find ideas or products that are not useful more creative (see Haaf & Rouder, 2017, 2018).

In order to examine the evaluation phase of the AUT and to explore individual differences in implicit conceptions of creativity, we apply the commonly used two-choice response time (RT) paradigm from the decision making literature (e.g., Krypotos, Beckers, Kindt, & Wagenmakers, 2015; Ratcliff, 1978; Wagenmakers, Ratcliff, Gomez, & McKoon, 2008), here in the form of the Creative-or-Not task (CON-task; Stevenson, Baas, & van der Maas, 2021). Rather than deciding on the correct response (e.g., Is CFREE a word? Yes or No?; Meyer & Schvaneveldt, 1971), on the CON-task people decide whether an AUT response is creative or not. For example, given the object “book” one must decide whether using it as a “roof tile” is creative (or not). The CON-task stimuli vary on the two dimensions of originality and utility, allowing us to unobtrusively assess individuals’ implicit values of originality and utility when judging creativity in the CON-task.

To gain insight into the cognitive basis of the evaluation process, we model CON-task data using the drift diffusion model (DDM; Ratcliff, 1978; Ratcliff & McKoon, 2008; Wabersich & Vandekerckhove, 2014). Over time, the individual accrues more and more information to answer the posed question and then ultimately crosses a decision boundary, for example to decide that “No, CFREE is not a word.” or to say “Yes, roof tile is a creative use for a Book” (see Figure 1). In the DDM the drift rate parameter reflects peoples’

tendency to respond “creative” or “not creative” based on their decisions as well as response times. By using the DDM to study creativity evaluation, we hope to disentangle otherwise potentially confounded, cognitive processes assumed to take place during the CON-task and to explore the role of response times. The originality and utility of the CON-task stimuli may not only influence decision outcomes, but also decision speed. For example, a highly original use may be easier to judge regarding creativity than a less original one and therefore generally be evaluated faster.




Figure 1. A graphical illustration of the DDM; α = boundary separation, indicating the evidence required to respond ‘Creative’ or ‘Not creative’; β = initial bias as a proportion of α indicating the tendency to choose one response over the other; δ = average drift rate, indicating the rate of information accumulation; τ = non-decision time, indicating the time used for processes apart from the actual decision-making.

As a plausibility check for our cognitive modeling approach, we also examine how well the relationship between individuals’ implicit values and their explicit opinions about the relevance of the two dimensions for creativity are aligned. Furthermore, we investigate whether those who value originality and utility also tend to produce more original and useful AUT responses, respectively. Regarding the latter, we build on Caroff and Besançon (2008)’s study who found a positive relationship for originality and called for research examining the same for “appropriateness” (i.e., a term often used as synonym for the utility dimension).

To this end, we conducted two studies. Study 1 is exploratory where we initially fit the

DDM model. In Study 2 we validate our pre-registered model from Study 1 and test specific hypotheses based on Study 1’s results.

General Method

Bayesian Hierarchical Diffusion Modeling of the CON-task

The DDM conceptualizes the response process in the CON-task as an interaction of several unobservable cognitive processes. Each of these is represented by a parameter (see Figure 1). We use the simplest complete version of the DDM comprising four parameters (Wabersich & Vandekerckhove, 2014). First, the model assumes that the decision whether a use is creative or not is initially determined by β . The β parameter stands for the a priori bias towards either of the decision options, so for the CON-task this is a preference for choosing “Yes, creative” or “No, not creative”. Second, the parameter δ stands for the drift rate. It reflects the tendency to respond “creative” or “not creative”. The higher the absolute drift rate, the faster the creativity evaluation, and the stronger the evidence for the decision. The DDM assumes that during each evaluation, noisy information is gradually accumulated. This accumulation ends when either of the two boundaries (i.e., “Yes, creative” or “No, not creative”) is reached. Third, the boundary separation parameter α reflects the distance between the two response boundaries and can be interpreted as response caution. For example, individuals who generally take longer to evaluate creativity, would have a greater boundary separation. Finally, the parameter τ refers to the non-decision time. This parameter aims to capture the processes taking place before and after the actual decision process whether something is creative or not, such as stimulus encoding and motor control processes. In sum, the DDM conceptualizes the process of judging whether a use is creative or not as a continuous random movement that begins somewhere between two boundaries and ends as soon as either of them is reached (Wabersich & Vandekerckhove, 2014).

In this paper, the most central DDM parameter is the drift rate. The assumption is

that stimulus originality and utility both positively affect the drift rate in that the more original and useful a CON-task stimulus is, the more positive the drift rate. Moreover, the drift rate is the only model parameter that we expect to be influenced by stimulus characteristics because the remaining parameters are already set before the decision of whether something is creative or not takes place (e.g., Vandekerckhove, Tuerlinckx, & Lee, 2011).

Bayesian hierarchical modeling. We estimated the model in a Bayesian hierarchical framework (Lee, 2011; Rouder & Lu, 2005; Vandekerckhove et al., 2011), allowing us to examine the data both at the population-level and at the individual-level. Hierarchical modeling provides rather conservative estimates of individual differences because it shrinks the individual effects towards the population mean (e.g., Efron & Morris, 1977; Haaf & Rouder, 2018).

We chose Bayesian estimation for two reasons. First, even without a hierarchical extension, applying the DDM to data is computationally expensive (e.g., Tuerlinckx, 2004). Extending it hierarchically makes the model quickly intractable when using the frequentist approach of maximum likelihood estimation (Vandekerckhove et al., 2011). Second, Bayesian inference has several advantages such as an intuitive treatment of uncertainty regarding the model parameters (Wagenmakers, 2009).

Model Specification

A detailed and complete model specification of the DDM used in Study 1 and 2 can be found in Appendix A. Here, we describe how we decomposed the drift rate parameter and the hierarchical structure of the model. To explore the influence of originality and utility when judging creativity in the CON-task, we regressed the drift rate on the originality and utility ratings of the stimuli. In both studies, we included random intercepts as well as random slopes to explore individual differences. Furthermore, because the response times

and proportions of “creative” responses vary considerably across the 64 CON-task stimuli, we also included random intercepts pertaining to the stimuli.

In both studies, we decomposed the drift rate as follows. Let $\delta_{(ij)}$ denote the drift rate for the i th participant, $i = 1, \dots, I$, in the j th trial or stimulus, $j = 1, \dots, 64$, of the CON-task, then

$$\delta_{(ij)} = \theta_{\delta(i)} + \phi_{\delta(j)} + \theta_{OR(i)}z_{OR} + \theta_{UT(i)}z_{UT}.$$

The parameters $\theta_{\delta(i)}$, $\theta_{OR(i)}$, and $\theta_{UT(i)}$ reflect participant i ’s drift rate decomposition. Specifically, $\theta_{\delta(i)}$ denotes the drift rate intercept, $\theta_{OR(i)}$ the originality effect, and $\theta_{UT(i)}$ the utility effect of individual i . Furthermore, $\phi_{\delta(j)}$ is stimulus j ’s individual deviation from the drift rate intercept. Lastly, z_{OR} and z_{UT} refer to z-scores of the originality and utility ratings of the stimuli.

For most of the remaining DDM parameters, we also incorporated random effects to examine individual differences. In particular, we allowed the boundary separation and the bias parameter to vary across individuals. However, in Study 1, we fixed the bias at the population level because we encountered identifiability issues when estimating random effects for β . We discuss these issues subsequently. Furthermore, in neither of the studies, we included random effects pertaining to the non-decision time parameter τ in the model. Instead, we constrained τ to be constant across participants because interpreting random effects for this parameter has shown to be problematic (Singmann, 2018).

To examine the interplay of the DDM parameters across participants, we also allowed the random effects pertaining to individuals to be correlated. As such, we assume that the individual effects are drawn from the same multivariate normal distribution with population means $[\mu_{\delta}, \mu_{\theta_{OR}}, \mu_{\theta_{UT}}, \mu_{\alpha}, \mu_{\beta}]^T$ and a variance-covariance matrix Σ , i.e.,

$$\begin{bmatrix} \theta_{\delta(i)} \\ \theta_{OR(i)} \\ \theta_{UT(i)} \\ \alpha_{(i)} \\ \beta_{(i)} \end{bmatrix} \sim \text{Multivariate-Normal} \left(\begin{bmatrix} \mu_{\delta} \\ \mu_{\theta_{OR}} \\ \mu_{\theta_{UT}} \\ \mu_{\alpha} \\ \mu_{\beta} \end{bmatrix}, \Sigma \right).$$

Σ allows for correlations across the random effects pertaining to the individuals. The random effects of the stimuli are orthogonal to the individual random effects. They are also assumed to be randomly sampled from a population distribution (of stimuli),

$$\phi_{\delta(j)} \sim \text{Normal}(0, \sigma_{\delta(j)}),$$

where 0 is the mean and $\sigma_{\delta(j)}$ is the standard deviation.

Since we estimated the model in the Bayesian framework, we needed to specify a prior distribution for each parameter. For Study 1, we used weakly informative priors that restricted the parameter space to a plausible range (see Appendix A). For Study 2, we used the insights gained from Study 1 and specified informative priors to test hypotheses in the Bayesian setting. We discuss these prior choices subsequently as well as in Appendix A.

Study 1

Data Collection Procedure and Materials

The data used for the DDM modeling in Study 1 was gathered as part of a joint data collection effort by the faculty of Psychology at the University of Amsterdam. Data collection took place over different sessions throughout 2016. The full sample consisted of 299 first year psychology students. The age range was 17.80- 41.35 years, ($M = 20.38$, $SD = 2.59$). Participants received course credit for their participation.

Alternative Uses Task. Participants completed a computerized version of the Alternative Uses Task (AUT; Guilford, 1967) used to assess their divergent thinking performance. The name of an object was presented on the screen, and participants had two minutes to type as many creative uses for the object as possible (e.g., the use “bath toy” for the object “brick”). Over the course of a year, a total of four objects were administered. Depending on the sessions the students participated in, they completed the task with a different object. Most participants completed the task for two objects, a few participants completed it only once. The objects for which participants were asked to generate uses were “brick”, “fork”, “towel”, and “paperclip”. Generated solutions were listed on the screen and new ones were continuously added. Two independent raters who were unaware of the research questions/hypotheses of this study separately scored participants’ answers with respect to originality and utility on a five-point scale (1 = not original/useful, 5 = very original/useful). Invalid responses were coded as 0. The interrater reliabilities (ICC) for the originality scores were .65, .68, .73 and .66 for “brick”, “fork”, “paperclip” and “towel” respectively. For utility scores the corresponding ICCs were .53, .60, .82 and .67 for “brick”, “fork”, “paperclip” and “towel” respectively.

Creative-or-not Task. Participants completed 64 trials of the CON-task. RTs as well as responses (“creative” or “not creative”) were recorded. Trials automatically counted as missing when participants did not answer within nine seconds. The stimuli used had been generated from a collection of AUT responses. Their originality and utility had been scored on a scale from 1 to 5 by two raters independently. The ICC was .81 for originality and .65 for utility. As stimulus ratings, we took the average originality and utility rating, respectively, across raters. The mean originality rating of the stimuli was $M = 2.98$ ($SD = 1.20$), and the mean utility rating was $M = 3.37$ ($SD = 1.07$). Stimulus originality and utility was negatively correlated, $r = -0.61$, 95% credible interval (CrI) $[-0.91, -0.45]$,

$$BF_{10} = 1.43 \times 10^{61}.$$

Importance Ratings of Originality and Utility. Participants also indicated, separately, how important they thought originality, utility, innovativeness, and appropriateness were when evaluating creativity (1 = not important at all to 5 = very important).

Results Study 1

The data cleaning is described in Appendix B. The cleaned dataset comprised 293 participants and 17806 trials. The mean RT across participants and trials after excluding data was 1.86 seconds ($Median = 1.61$, $SD = 0.92$). The overall proportion of “creative” responses was 53.83%. The mean RT for “creative” responses was 1.87 s ($SD = 0.93$), and 1.86 s ($SD = 0.92$) for “not creative” responses. Figure 2A shows that RTs for some stimuli were longer than for others. In particular, there was a small tendency for stimuli with higher originality and utility ratings to be answered more slowly. Furthermore, Figure 2B shows that on average, participants did respond similarly fast for both “not creative” and “creative” responses. It seems that there were no obvious effects of RT in the data. Nevertheless, we decided to proceed with the diffusion model as planned.

Model Fit. We fitted the model using the R package *brms* (Bürkner, 2018) which works with Stan to draw samples from the posterior distribution of Bayesian models (Carpenter et al., 2017). We ran 4 chains with 5000 iterations each. 1500 iterations per chain were used as warmup to adapt the sampler. Consequently, our analyses were based on a

¹ Here and for all subsequently reported correlations, we conducted Bayesian correlation analyses using the Bayes factor package including the default prior scale (Morey & Rouder, 2018). Specifically, we used Bayes factors to quantify the evidence for a correlation ($H_1 : \rho \neq 0$) as opposed to no correlation ($H_0 : \rho = 0$;) and report it together with the posterior mean of the correlation coefficient and the corresponding credible interval.

total of 14000 iterations.²

We decided to not allow the bias parameter β to vary across individuals because when we did, the random effects of the bias and drift rate intercept were highly correlated, suggesting identifiability issues. We therefore estimated the DDM with only a population-level bias parameter. The remaining DDM parameters were not affected by this change.

We performed several model diagnostics procedures and inspected the model fit. There were no signs of non-convergence, with 0 divergent transitions and \hat{R} values (Gelman & Rubin, 1992) below 1.01 (Vehtari, Gelman, Simpson, Carpenter, & Bürkner, 2020). Additionally, we assessed the model fit using posterior predictive checks (see the online supplementary material). Overall, there was no major model misfit, with the occasional misfit at the individual level.

Modeling Results. A summary of the estimated fixed and random effects parameters can be found in Table 1 and 2, respectively, and a summary table with the random effects correlation parameters in Appendix B.

Regarding the fixed effects, our main focus of interest was on $\mu_{\theta_{OR}}$ and $\mu_{\theta_{UT}}$, the

² For all analyses, we used R (Version 3.6.1; R Core Team, 2019) and the R-packages *BayesFactor* (Version 0.9.12.4.2; Morey & Rouder, 2018), *bayesplot* (Version 1.7.0; Gabry, Simpson, Vehtari, Betancourt, & Gelman, 2019), *brms* (Version 2.10.0; Bürkner, 2017), *coda* (Version 0.19.3; Plummer, Best, Cowles, & Vines, 2006), *corrplot2017* (Wei & Simko, 2017), *cowplot* (Version 1.0.0; Wilke, 2019), *DescTools* (Version 0.99.29; al., 2019), *dplyr* (Version 1.0.0; Wickham, François, Henry, & Müller, 2020), *ggplot2* (Version 3.3.3; Wickham, 2016), *gridExtra* (Version 2.3; Auguie, 2017), *here* (Version 0.1; Müller, 2017), *magrittr* (Version 2.0.1; Bache & Wickham, 2020), *Matrix* (Version 1.2.17; Bates & Maechler, 2019), *msm* (Version 1.6.7; Jackson, 2011), *mvtnorm* (Version 1.0.11; Genz & Bretz, 2009), *papaja* (Version 0.1.0.9997; Aust & Barth, 2020), *psych* (Version 1.8.12; Revelle, 2018), *Rcpp* (Version 1.0.2; Eddelbuettel & François, 2011; Eddelbuettel & Balamuta, 2017), *rstan* (Version 2.19.2; Stan Development Team, 2019a), *StanHeaders* (Version 2.19.0; Stan Development Team, 2019b), *stringr* (Version 1.4.0; Wickham, 2019), *tibble* (Version 3.0.4; Müller & Wickham, 2020), and *tidyr* (Version 1.0.0; Wickham & Henry, 2019).

overall effects of stimulus originality and utility on the drift rate δ . For $\mu_{\theta_{OR}}$, the posterior mean was 0.41, and the 95% CrI was [0.32, 0.51]. For $\mu_{\theta_{UT}}$, the posterior mean was 0.10, and the 95% CrI was [0.01, 0.19]. Both posterior means were positive. However, while the 95% credible interval (CrI) of $\mu_{\theta_{OR}}$ did not include zero, the CrI of $\mu_{\theta_{UT}}$ was very close to zero.

In general, all estimated posterior means of the remaining fixed effects parameters seem plausible as the CrIs were rather narrow and the parameters lie within a reasonable range (see Table 1). On average, there was no a-priori bias towards the response options “creative” or “not creative” (see μ_{β} in Table 1). This suggests that before stimulus onset, participants were on average equally likely to choose either of the two response options. However, the boundary separation, or response caution, was higher than found in most applications of the diffusion model (e.g., Matzke & Wagenmakers, 2009, see μ_{α} in Table 1). One explanation for this rather high value is that the RTs in the CON-task were considerably slower than RTs in tasks typically modeled by the diffusion model.

We found substantial individual differences in all variability parameters. Notably, the results showed substantial variability across participants in the originality and utility effects on the drift rate. The posterior means of σ_{OR} and σ_{UT} were 0.23 and 0.19, respectively. The posterior means and credible intervals of all variability parameters are listed in Table 2. None of the CrIs included zero suggesting considerable variability across stimuli and across participants.

Figures 3A and B visualize this variability in the originality and utility effects on the drift rate by depicting the posterior means of the individual originality slopes $\theta_{OR(i)}$ and utility slopes $\theta_{UT(i)}$ and their corresponding CrIs in increasing order. The figures show substantial individual differences. There were even 11 participants with a negative posterior mean (3.75%). However, the 95%CrI of these estimates included zero. In total, the CrI of 51 individuals included zero (17.41%), suggesting that, at the very least, for some individuals, the effect of stimulus originality on the drift rate was weaker and for some it was stronger.

Participants further differed in their utility slopes. As shown in Figure 3B, for some,

the utility effect on the drift was around zero, for some it was positive, and for a few the effect was even negative. Specifically, the majority of the individual CrIs ($n = 208$; 70.99%) included zero, 79 (26.96%) of them excluded and were above zero, and 6 (2.05%) individual CrIs excluded and were below zero.

Individual differences also manifested themselves in a negative correlation between the originality and utility slopes. Here the posterior mean of $\rho_{\sigma_{\theta_{OR}}\sigma_{\theta_{UT}}}$ was -0.44, 95%CrI [-0.60, -0.27], and the correlation between the individual originality and utility slopes, based on the posterior means, was $r = -0.67$, 95%CrI [-0.73, -0.60], $BF_{10} = 6.03 \times 10^{37}$. It is expected that the former is greater than the latter because $\rho_{\sigma_{\theta_{OR}}\sigma_{\theta_{UT}}}$ is a population parameter taking uncertainty into account and r reflects the data in our sample. The greater the individual effect of stimulus originality on the drift rate, the smaller the effect of stimulus utility and vice versa. This correlation could also be explained by the substantial negative correlation between the originality and utility ratings for the CON-task stimuli ($r = -0.60$). We ruled this out by re-estimating the DDM when excluding the stimuli that contributed the strongest to the negative correlation. Specifically, we excluded the data from 20 items leaving us with CON-task data based on 44 stimuli. Excluding those items reduced the stimulus originality-utility ratings correlation from $r = -0.60$ to $r = -.15$. Despite this reduced correlation across stimuli, the negative correlation across originality and utility effects on individual drift rates remained substantial, $r = -.49$, 95%CrI [-0.67, -0.29] (see online supplementary). This result suggests that the correlation across effects is not (solely) a function of stimulus characteristics.

Figure 4A, shows the multivariate, joint posterior distribution of the originality and utility effects and Figure 4B the individual posterior means and corresponding standard deviations to visualize this correlation. An overview table of the correlations among all random effects parameters can be found in Table 3 in Appendix C. Figure 4B further depicts two clusters that are the results from a k-means cluster analysis: it seems that one cluster comprises individuals with a positive effect of utility and a rather small (or even negative)

effect of stimulus originality, and the other one individuals with a stimulus utility effect around zero and a positive effect of originality.

CON-task and self-report ratings of originality and utility. As a plausibility check for our rationale behind the drift rate regression, we examined whether the self-reported importance ratings of originality, innovativeness, utility, and appropriateness corresponded to the originality and utility effects on the drift rate. We summed up the ratings of appropriateness and utility and innovativeness and originality, respectively. There was a positive correlation between participants' importance ratings of originality and the posterior means of their originality slopes, $r = 0.32$, 95%CrI [0.21, 0.42], $BF_{10} = 9.76 \times 10^5$ and between their ratings of utility and utility slopes, $r = 0.36$, 95%CrI [0.26, 0.46] , $BF_{10} = 1.68 \times 10^8$. The more participants indicated that originality was important when determining whether something is creative or not, the greater their influence of stimulus originality on their drift rate. The more they indicated that utility was important, the greater their effect of stimulus utility on their drift rate. There were also negative correlations between the originality importance ratings and the utility slopes, $r = -0.33$, 95%CrI [-0.43, -0.23] , $BF_{10} = 4.67 \times 10^6$, and between the utility ratings and the originality slopes, $r = -0.36$, 95%CrI [-0.46, -0.26], $BF_{10} = 1.74 \times 10^8$.

CON-task and divergent thinking. Given the substantial variability in the stimulus originality and utility effects on the drift rate, we explored whether this variability was related to variability in divergent thinking performance as assessed by the AUT. To this end, we computed correlations among the individual posterior means and participants' AUT performance scores. The data cleaning for the AUT task is described in Appendix B.

We found a positive correlation between AUT originality scores and the originality slope posterior means, $r = 0.17$, 95%CrI [0.05, 0.30] , $BF_{10} = 6.53$, suggesting that the more original the AUT responses, the greater the influence of originality on the drift rate. We also found a negative correlation between the posterior means of the utility slopes and

participants originality scores, $r = -0.15$, 95%CrI [-0.27, -0.03], $BF_{10} = 2.66$. The more original responses participants produced in the AUT, the smaller their effects of stimulus utility on the drift rate. However, there was no correlation between the AUT utility scores and the stimulus utility effects on the drift rate, $r = 0.01$, 95%CrI [-0.12, 0.13], $BF_{01} = 6.56$, and also no correlation between AUT utility scores and stimulus originality effects, $r = -0.01$, 95%CrI [-0.13, 0.12], $BF_{01} = 6.60$.

Since the application of the DDM to creativity is novel and since Study 1 was conducted in a rather exploratory manner, we aimed to assess the robustness of our findings in a second, preregistered Study (<https://osf.io/7gt45/>). In Study 2, we therefore specified hypotheses based on Study 1's results as well as previous research and re-fitted the DDM on an independent dataset.

Study 2

Based on what we had learned from Study 1 and based on previous research, we expected positive effects of stimulus originality (H1) and stimulus utility (H2) on the drift rate. We also expected the effect of stimulus originality to be larger than the effect of stimulus utility (H3). Furthermore, given the observed substantial negative correlation between the individual stimulus originality and utility effects in Study 1, we expected a negative correlation among those effects in Study 2 (H4). Specifically, we hypothesized that the greater an individuals' effect of stimulus originality on the drift rate, the smaller the effect of stimulus utility would be and vice versa. Since in Study 1, we observed substantial individual differences in the extent to which stimulus originality and utility influenced the drift rate, we also expected non-zero variability across individuals in those effects (H5a and H6a).

Previous research suggests that originality plays a superior role in creativity judgments compared to utility (e.g., Caroff & Besançon, 2008; Diedrich et al., 2015; Runco & Charles, 1993). Based on these findings we tested the hypothesis that everyone would have a positive

effect of stimulus originality on the drift rate (H5b; Haaf & Rouder, 2017, 2018). Note that this hypothesis deviated from Study 1's findings where some people even seemed to have a negative effect. Regarding the effects of stimulus utility on the drift rate, we expected that some individuals would have a positive effect, some a negative effect, and some no effect (H6b). Given Study 1's results, we also expected that individuals' stimulus originality and utility effects would be positively associated with their self-reported importance ratings of originality and utility for creativity. More specifically, we expected individual originality effects on the drift rate would increase as the self-reported importance ratings of originality increase (H7a), and individual utility effects on the drift rate would increase as the importance ratings of utility increase (H7b).

Finally, we specified hypotheses regarding the association between CON-task judgements and AUT performance. We expected that originality scores on the AUT would be positively correlated with stimulus originality effects and negatively correlated with stimulus utility effects on drift rates (H8a and H8b). Similarly, we expected that AUT utility scores would be positively correlated with stimulus utility effects on CON-task drift rates and negatively correlated with stimulus originality effects (H9a and H9b). Note that Study 1 did not support H9a and H9b. However, in both cases, we did not consider the evidence for the null to be convincing. We still expected that the more useful one's AUT responses are, the more one values utility when judging creativity and the more one disregards originality. Study 2 served as a robustness check for this belief.

Data Collection Procedure and Materials

Data collection was again centrally organized by the faculty of Psychology at the University of Amsterdam and took place over different sessions throughout 2019. All tasks related to creativity were administered during the same session in the order listed below. In total, 176 first-year psychology students completed the CON-task, the age range was 17.87-47.04 years, ($M = 20.50$, $SD = 2.98$, and participants again received course credit for their

participation.

Alternative-Uses-Task Participants completed the AUT (Guilford, 1967) for the objects “brick” and “paperclip” and were again given two minutes for each object. Two independent raters who were unaware of the research questions/hypotheses of this study again separately scored participants’ answers with respect to originality and utility on a five-point scale and coded invalid responses as zero. The interrater agreement can be considered good: for the object “brick”, the ICC was 0.67 for originality and 0.64 for utility. For the object ‘paperclip the ICC was 0.73 for originality and 0.71 for utility.

Creative-Or-Not task Participants again completed the same 64 stimuli as the participants in Study 1.

Self-reported importance-ratings On four separate items, participants again indicated how important they thought originality, innovativeness, appropriateness, and utility were when deciding whether something is creative or not.

Results Study 2

The hypotheses and analysis plan for Study 2 were preregistered before seeing the data. We employed the same exclusion criteria as in Study 1 (see Appendix B).

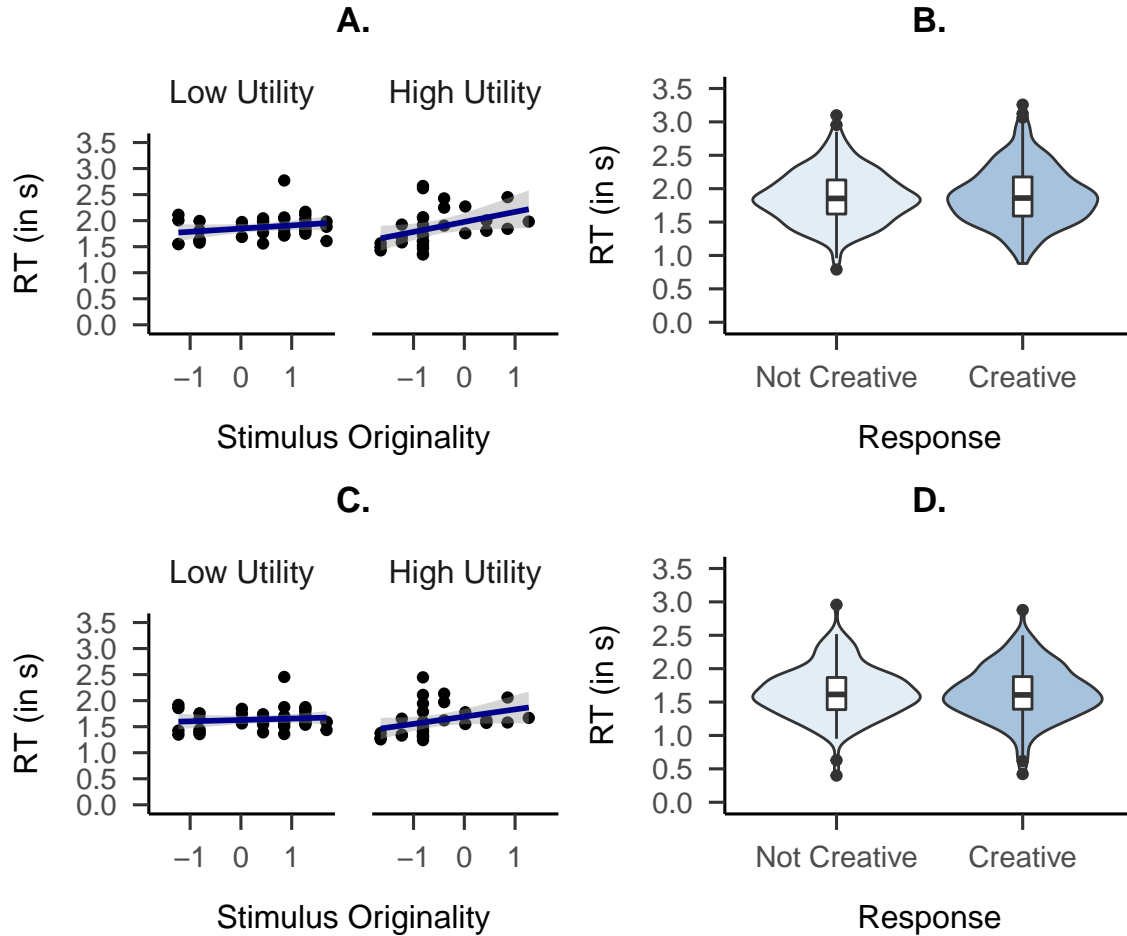


Figure 2. A and C show the mean stimulus response times as a function of stimulus originality and utility in Study 1 and 2, respectively. Each dot represents a stimulus. Low and high utility stimuli were categorized by median split. B and D show the distribution of mean response times for 'creative' and 'not creative' responses. Each dot represents a participant.

The descriptives were very similar to Study 1. The mean RT across participants and trials was 1.62 s and the median 1.44 s ($SD = 0.79$). The overall proportion of “creative” responses was 0.57. The mean RT for “creative” responses was 1.62 s ($SD = 0.78$). For “not creative” responses, this was 1.63 s ($SD = 0.79$). Stimuli with higher scores on originality and utility again appeared to be answered more slowly, and RTs for two responses were very similar (see Figures 2C and D).

Model Fit. In contrast to Study 1, we used informative priors for the stimulus originality and utility effects on the drift rate, based on Study 1’s estimation results. Specifically, we specified truncated normal distributions as priors for $\mu_{\theta_{OR}}$, and $\mu_{\theta_{UT}}$,

$$\mu_{\theta_{OR}}, \mu_{\theta_{UT}} \sim \text{Normal}^+(0, 0.2),$$

All remaining priors were the same as in Study 1 except for β (see Appendix A). We again fitted the model using the R package *brms* (Bürkner, 2018), ran 4 chains with 4000 iterations of which 500 iterations per chain were used as warmup, leaving us with 14000 iterations to base the analysis on.

We again inspected the model diagnostics and model fit. There were no signs of non-convergence, with 0 divergent transitions, and \hat{R} values below 1.01 (Vehtari et al., 2020). Moreover, we again assessed the model fit using posterior predictive checks. The fit assessment was favorable (see the online supplementary material). As a robustness check, we also re-estimated the model including participants that we excluded based on exclusion criterion 3 (fewer than 47 remaining trials; see the online supplementary materials).

Modeling Results. Summary statistics of the estimated model parameters are shown in Table 1 and 2. A table with the correlations among the random effects can be found in Appendix B. Overall, the estimated DDM parameters were very similar to the ones in Study 1.

Table 1

Posterior mean, standard deviation of the posterior distribution, 95% credible interval and \hat{R} statistic for the fixed effects (population-level) parameters

	Study 1				Study 2			
	Mean	SD	LB	UB	Mean	SD	LB	UB
μ_δ	0.06	0.04	-0.01	0.14	0.15	0.05	0.05	0.25
μ_α	2.96	0.03	2.91	3.01	2.72	0.04	2.64	2.80
μ_β	0.50	0.00	0.49	0.50	0.49	0.01	0.48	0.50
$\mu_{\theta_{OR}}$	0.41	0.05	0.32	0.51	0.40	0.05	0.30	0.51
$\mu_{\theta_{UT}}$	0.10	0.05	0.01	0.19	0.10	0.05	0.01	0.20
τ	0.28	0.00	0.27	0.28	0.29	0.00	0.29	0.29

Note. $\mu_{\theta_{OR}}$, $\mu_{\theta_{UT}}$, and μ_δ are standardized estimates as the originality and utility ratings of the stimuli are z-scores. SD = standard deviation; LB = lower bound; UB = upper bound.

Table 2

Posterior mean, standard deviation of the posterior distribution, 95% credible interval and \hat{R} statistic for the variability parameters. σ_{δ_ϕ} denotes the variability across stimuli

	Study 1				Study 2			
	Mean	SD	LB	UB	Mean	SD	LB	UB
σ_{δ_ϕ}	0.27	0.03	0.23	0.33	0.33	0.03	0.27	0.39
σ_{δ_v}	0.27	0.01	0.25	0.3	0.36	0.03	0.32	0.42
σ_{OR}	0.23	0.01	0.21	0.26	0.20	0.02	0.16	0.24
σ_{UT}	0.19	0.01	0.17	0.22	0.20	0.02	0.17	0.24
σ_α	0.4	0.02	0.36	0.44	0.46	0.03	0.41	0.52
σ_β	-	-	-	-	0.05	0.01	0.04	0.06

Note. SD = standard deviation; LB = lower bound; UB = upper bound.

Hypotheses testing. To test our hypotheses, we used the Savage-Dickey method (Dickey, 1971) of approximating Bayes factors. In this method, the Bayes factor is computed by a ratio of the prior and posterior density at the value zero. Assessing H1, we computed a Bayes factor comparing how well the hypothesis of a positive effect of stimulus originality on the drift rate predicted the data in comparison to the null hypothesis. All posterior samples were greater than zero. Therefore, the evidence in favor of H1 can be regarded as greater than 14000. In contrast, there was mere anecdotal evidence for an effect of utility on the drift rate, the Bayes factor was 2.33. This means that the data was 2.33 times more likely to have occurred under H2 than under the null hypothesis. Not surprisingly, the Bayes factor for H3, that the overall effect of stimulus originality is greater than the effect of stimulus utility

(rather than the other way around), was again greater than 14000.

We examined the correlation between the stimulus originality and utility effects on the drift rate (H4). The random effects correlation as estimated by the model, $\rho_{\sigma_{\theta_{OR}}\sigma_{\theta_{UT}}}$, had a posterior mean of -0.34, 95%CrI $[-0.58, -0.09]$, and the correlation based on the individual posterior means was $r = -0.64$ 95%CrI $[-0.73, -0.54]$, $BF_{10} = 1.45 \times 10^{16}$. We again explored the robustness of the correlation between the originality and utility effects on the drift rate and by estimating the DDM based on data from a subset of weakly correlated stimuli (see Study 1)³. Unlike in Study 1, the correlation between the effects became considerably smaller. The posterior mean of the correlation was reduced to the size of the stimulus correlation in the reduced item pool ($r = -.17$), suggesting that the negative correlation between originality and utility stimulus effects on the drift rate was not as robust as in Study 1.

The individual differences were again reflected in the variability parameters of the originality and utility effects on the drift rate. The posterior mean of the variability parameters for σ_{OR} was 0.20 and for σ_{UT} this was also 0.20. The Savage-Dickey density ratio cannot be used to compute a Bayes factor quantifying the evidence for individual variability. However, the 95% credible intervals of both variability parameters did not include zero, supporting H5a and H6a that the variability was substantial. To further examine the extent of individual variation, we again plotted individual posterior means and credible intervals in increasing order (see Figures 3C and D). Regarding originality, the CrI of 24 individuals included zero (15.79%) indicating that some individuals' drift rates were only weakly (if at all) determined by the stimuli's originality. Out of these participants, there were even 3 participants with a negative posterior mean (1.97%). Participants also again differed in their utility slopes. The majority of the individual CrIs ($n = 118$; 77.63%) included zero, 32 (21.05%) of them excluded and were above zero, and 2 (1.32%) individual CrIs excluded and were below zero.

³ This analysis was not pre-registered.




Figure 3. The plots show the posterior means and the 95 % credible interval (CrI) for each participant in increasing order. The dashed horizontal line denotes the population-level posterior means μ_{OR} , μ_{UT} . CrIs colored in red included zero. Plot A and C show the individual estimates of the originality effect in Study 1 and Study 2, respectively. Only a few lower bounds of the CrIs were below zero. Plot B and D show the individual estimates of the utility effect in Study 1 and Study 2, respectively. Some CrIs are above zero, some around zero, and a few below zero.

In line with Haaf and Rouder (2017) and Haaf and Rouder (2018), we tested the hypothesis that all individuals have a positive effect of stimulus originality on the drift rate using the encompassing prior approach (H5b; Klugkist, Laudy, & Hoijtink, 2005; Klugkist & Hoijtink, 2007). Here, we compared the predictive accuracy of the hypothesis that everyone's originality slope is positive to the hypothesis that originality effects can be positive, zero or

negative. The constraint that everyone has a positive effect was not fulfilled in any of the iterations. Therefore, the Bayes factor for the latter, unconstrained hypothesis can be considered at least 2,230.16, (assuming that the next iteration might be the first in favor of the ordinal constraint) suggesting that not everyone has a positive originality effect on the drift rate. Regarding the individual utility effects, we tested the hypothesis that some individuals have a positive effect of stimulus utility on the drift rate, some have a negative effect, and some have no effect (H6b). Here, the data provided overwhelming support for the unconstrained hypothesis H6b over the hypothesis that everyone has a positive utility effect as the Bayes factor was again at least 2,230.16 (assuming that the next iteration might be the first in favor of the ordinal constraint).

CON-task and self-report ratings of originality and utility. To test the hypothesis that stimulus originality and utility effects on the drift rate were positively correlated with self-reported importance ratings of originality and utility, respectively (H7a and b), we conducted one-sided Bayesian correlation tests again using sum scores. There was a positive correlation between the originality sum scores and the individual originality effects, $r = 0.17$, 95%CrI [0.03, 0.32], $BF_{10} = 2.54$ as well as a positive correlation between the utility sum scores and the individual utility effects, $r = 0.24$, 95%CrI [0.08, 0.38], $BF_{10} = 24.98$, supporting both H7a and H7b. We again conducted a cluster analysis on the individual posterior means (see Figure 4), which yielded practically the same result as in Study 1⁴.

⁴ This analysis was not pre-registered.

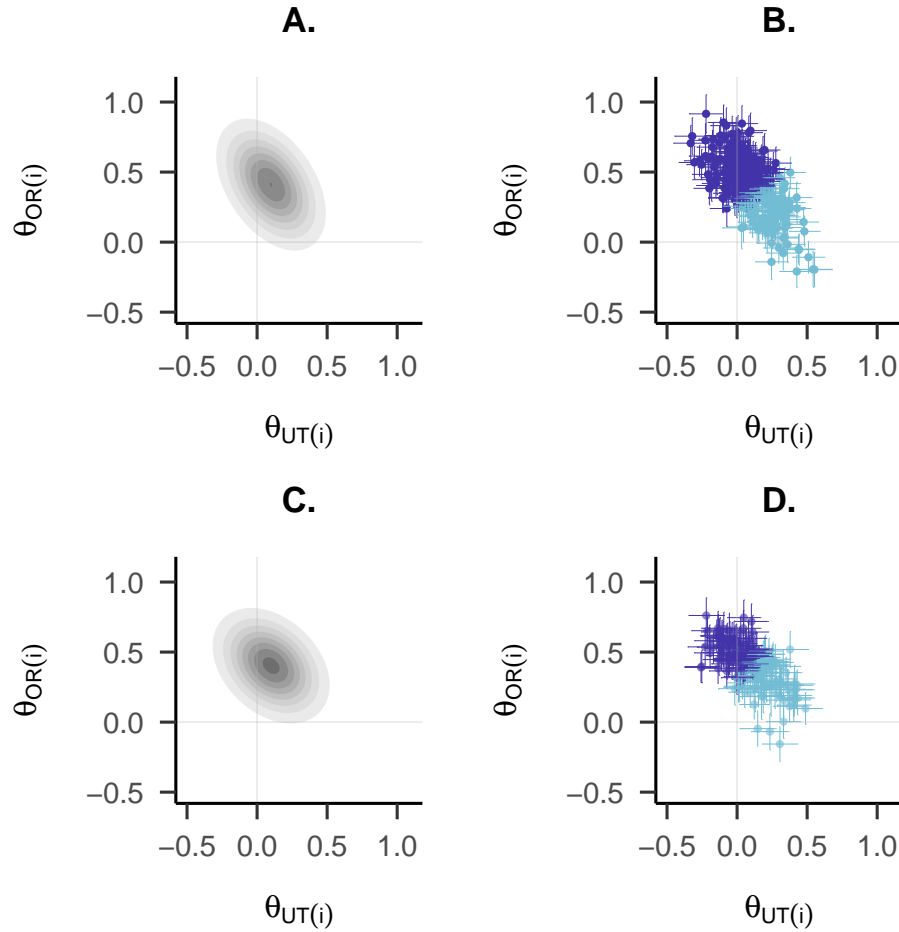


Figure 4. A and C show the multivariate, joint posterior distribution of the originality effects ($\theta_{OR(i)}$) and utility effects ($\theta_{UT(i)}$) on the drift rate, in Study 1 and Study 2, respectively. Darker areas indicate greater density. B and D show the individual posterior means of the originality and utility effects including the standard deviations. Each dot represents a participant. The dots are colored according to the results of k-means cluster analysis. Both plots depict the negative correlation between the originality and utility effects.

CON-task and divergent thinking. To examine associations between the stimulus originality and utility effects on the drift rate and creative performance (H8 and H9), we followed the same data cleaning procedures as described in Study 1. Again, few participants ($n = 12$) did not complete the task. Since we had directed hypotheses, we used one-sided correlation tests. There was a positive correlation between the stimulus originality

effects and the AUT originality scores, $r = 0.24$, 95%CrI [0.08, 0.39], $BF_{10} = 21.92$. The data also showed that the AUT originality scores and the stimulus utility effects from the CON-task were negatively correlated, $r = -0.17$, 95%CrI [-0.32, -0.03], $BF_{10} = 2.19$ (H8b). Furthermore, there was a positive correlation between the AUT utility scores and the stimulus utility effects, $r = 0.20$, 95%CrI [0.05, 0.35], $BF_{10} = 5.15$ (H9a), and a negative correlation between the AUT utility scores and the stimulus originality effects on the drift rate, $r = -0.25$, 95%CrI [-0.40, -0.09], $BF_{10} = 27.58$ (H9b).

Discussion

In this paper, we aimed to study the cognitive basis of how people decide whether an idea is creative or not during divergent thinking. Divergent thinking is a two-phase process where people come up with ideas (ideation phase) and decide whether or not to use them (evaluation phase; Basadur, 1995; Guilford, 1967; Runco & Basadur, 1993). We focused on the evaluation phase and used Bayesian cognitive modeling to better understand it. Using a novel application of the drift diffusion model (DDM) we aimed to understand how people weigh the two components of creativity, originality and utility, when deciding whether an idea is creative or not. We used the Creative or Not (CON) task, a timed two-choice decision-making task, where people are presented with ideas of how to use an everyday object (e.g., use a book as a roof tile) and then decide whether these are creative or not. The DDM drift rate parameter represents people's tendency to respond "creative" or "not creative" based on their judgements and response times. The CON-task stimuli varied on how original and how useful they were (e.g., using a book to read is useful, but not original, whereas using a book as roof tile can be both). This way we could estimate individual differences in how originality and utility implicitly contributed to people's creativity judgements. Our main findings were: (1) stimulus originality strongly influenced participants' decisions, whereas stimulus utility was only somewhat related to their tendency towards creative responses; (2) there were large individual differences in how much participants took the originality and utility of presented ideas into account; (3) participants' implicit values of originality and utility on the CON-task were aligned with their self-reported importance ratings of originality and utility for creativity; (4) Alternative Uses task (AUT) scores for originality and utility coincided with how much the originality and utility of the presented CON-task idea weighed in on their decisions. We now discuss each of these findings in more detail.

Our finding that people value originality when evaluating creativity, evidenced by a substantial overall effect of originality on DDM drift rates, are in line with previous studies

showing that people associate originality more so than utility with creativity (e.g., Caroff & Besançon, 2008; Diedrich et al., 2015; Runco & Charles, 1993). In both studies, the utility of ideas was only somewhat related to participants' tendency towards creative responses. In earlier work we also see that utility, or one of its creative synonyms (e.g., appropriateness, value) is less valued when it comes to judging creativity (e.g., Diedrich et al., 2015). Future work with the CON-task could include additional information about presented ideas, such as “surprise” (Boden, 2004) and “impact” (Sternberg & Lubart, 1996), to examine how these components implicitly influence people's creativity judgments.

Interestingly, just as researchers sometimes disagree on the definition of creativity (Simonton, 2018), the homogenous group of psychology students in our study also varied greatly in their conceptions of creativity. We found both quantitative and qualitative individual differences in how they valued originality and utility (Haaf & Rouder, 2017, 2018). Originality was usually important, but a few people even seemed to regard original stimuli to be “not creative”. Regarding utility, there were large individual differences too. Some participants seemed to value utility, a few rated any useful idea as “not creative”, and some appeared to disregard it entirely. Future research should look to generalize these results to other populations or types of creativity judgments. For example, adolescents appear to place more value on originality and less on utility (Stevenson et al., 2021), so there may be a developmental trend. Also, judging the creativity of AUT ideas is interesting as the task is used so often in psychological assessment. But, would utility be considered more valuable in real-world situations, e.g., judging the creativity of ideas to combat climate change?

After the CON-task, we asked participants to rate how important they thought originality and utility were for creativity. These self-reported relevance ratings more or less corresponded to their implicit values of originality and utility on the CON-task. So, the more important someone reported originality or utility was for creativity, the more they took it into account on the CON-task. So it appears that the CON-task provides an indirect way to tap into what people consider creative. Similar to Caroff and Besançon (2008), our results

also showed that the more participants took originality and utility into account on the CON-task, the more original and useful their responses to the AUT were. We cannot talk about causality, but it seems plausible that people’s implicit values of originality and utility influence what ideas they produce and how they evaluate and select ideas during divergent thinking, like on the AUT.

Methodological Implications

To our knowledge, this is the first time a mathematical model such as the DDM has been applied to the evaluation phase of divergent thinking. The DDM is generally applied to timed two-choice decision tasks where there is clearly a correct answer, such as lexical decision tasks (e.g., Wagenmakers et al., 2008). Therefore, applying the DDM to the CON-task, where there is no correct answer and we are basically assessing what factors influence people’s tendency to decide if an idea is creative, is novel. Applying the DDM generally worked well on both CON-task studies, as suggested by the good model fit. However, the drift rate seemed to be mainly driven by decisions rather than response times (RTs) and RTs did not vary much across stimulus originality and utility ratings. At best, there was a tendency for stimuli with higher originality and utility ratings to be answered more slowly. A possible explanation is that the effects of originality and utility on the drift rate (i.e., people’s tendency towards deciding “creative”) were negatively correlated in both studies. This means that the more participants valued originality, the less they valued utility and vice versa. There appears to be a trade-off where people balance one dimension against the other when judging creativity. This may explain why stimuli with both higher originality and utility took a little longer to process.

However, the lack of influence of RTs begs the question of whether the DDM is not overly complex for CON-task data. For example, a hierarchical probit analysis⁵ of our data resulted in similar conclusions. It would be interesting to examine whether RT effects are

⁵ exploratory, not pre-registered, see online supplementary materials

more apparent if participants are given less time (e.g., 3s versus the 9s in this study); especially since in the real-world people sometimes need to make split-second choices about which idea to pursue. Also, instructions that emphasize speed could affect the evaluation process, just as those of quantity or quality affect the ideation phase of divergent thinking (Said-Metwaly, Fernández-Castilla, Kyndt, & Van den Noortgate, 2020).

Conclusion

This paper demonstrates a new approach to studying the evaluation process during divergent thinking. Our novel application of the drift diffusion model provides a mathematical method to study how people decide what's creative. The results imply that it is important to take both originality and utility into account when examining how people judge creativity. Also, given how conceptions of creativity vary, our findings suggest that when assessing creativity using divergent thinking tasks like the AUT, managers or researchers should clarify what the rating criteria are to provide a fair playing field for all.

Appendices

Appendix A

This Appendix contains the DDM specification. Model differences between the studies are mentioned where necessary.

Let $\mathbf{Y}_{(ij)}$ denote a response vector of the decision and response time $(X_{(ij)}, T_{(ij)})$ for the i th participant, $i = 1, \dots, I$ in the j th trial, $j = 1, \dots, J$ of the CON task. The bivariate data $\mathbf{Y}_{(ij)}$ is assumed to be distributed according to a Wiener distribution,

$$\mathbf{Y}_{(ij)} \sim \text{Wiener}(\alpha_{(ij)}, \beta_{(ij)}, \tau_{(ij)}, \delta_{(ij)}),$$

with the four model parameters boundary separation α , bias β , non-decision-time τ , and drift rate δ . The Wiener distribution is a joint density function of deciding whether a use is creative or not, $X_{(ij)}$, at time $T_{(ij)}$ (Vandekerckhove et al., 2011).

The double index notation suggests that, in principle, the four parameters may differ across participants, as well as across trials. To reduce model complexity, we constrain the model in several ways. First, we treat all parameters constant across trials. Second, at the participant level, we allow participants to differ in three out of four parameters, α , β , and δ . The non-decision time parameter τ be constrained to be constant across trials as well as across participants, i.e., $\tau_{(ij)} = \tau$, because interpreting random effects for the non-decision time parameter has shown to be problematic (Singmann, 2018). We treat differences across individuals as random effects, assuming that participants are a sample from a population distribution.

Our main focus of interest be on the drift rate parameter δ . It is the only model parameter assumed to be influenced by stimulus characteristics because the remaining

parameters are already set before the decision-making whether something is creative or not takes place (e.g., Vandekerckhove et al., 2011). To examine the influence of originality and utility when evaluating creativity, we regress δ on the originality and utility ratings of the stimuli. We include random intercepts as well as random slopes to take interindividual variation into account. Furthermore, because the response times and proportions of “creative” responses have shown to vary considerably across the 64 task stimuli, we also include random intercepts pertaining to the stimuli. Mathematically, we express the above described as follows,

$$\delta_{(ij)} = \theta_{\delta(i)} + \phi_{\delta(j)} + \theta_{OR(i)}z_{OR} + \theta_{UT(i)}z_{UT}.$$

The parameters $\theta_{\delta(i)}$, $\theta_{OR(i)}$, and $\theta_{UT(i)}$ reflect participant i 's drift rate decomposition. Specifically, $\theta_{\delta(i)}$ denotes the drift rate intercept, $\theta_{OR(i)}$ the originality effect, and $\theta_{UT(i)}$ the utility effect of individual i . Furthermore, $\phi_{\delta(j)}$ is stimulus j 's individual deviation from the drift rate intercept. Lastly, z_{OR} and z_{UT} refer to z-scores of the originality and utility ratings of the stimuli.

In Study 1, the boundary separation parameter has random and fixed effects and the bias parameter is fixed. In Study 2, both the boundary separation and the bias parameter are allowed to vary across individuals. All random effects pertaining to the individuals to be correlated in both studies. As such, we assume that the random effects are drawn from the same multivariate normal distribution with variance-covariance matrix Σ , i.e.,

$$\begin{bmatrix} \theta_{\delta(i)} \\ \theta_{OR(i)} \\ \theta_{UT(i)} \\ \alpha_{(i)} \end{bmatrix} \sim \text{Multivariate-Normal} \left(\begin{bmatrix} \mu_{\delta} \\ \mu_{\theta_{OR}} \\ \mu_{\theta_{UT}} \\ \mu_{\alpha} \end{bmatrix}, \Sigma \right),$$

in Study 1, and

$$\begin{bmatrix} \theta_{\delta(i)} \\ \theta_{OR(i)} \\ \theta_{UT(i)} \\ \alpha_{(i)} \\ \beta_{(i)} \end{bmatrix} \sim \text{Multivariate-Normal} \left(\begin{bmatrix} \mu_{\delta} \\ \mu_{\theta_{OR}} \\ \mu_{\theta_{UT}} \\ \mu_{\alpha} \\ \mu_{\beta} \end{bmatrix}, \mathbf{\Sigma} \right),$$

in Study 2.

$\mathbf{\Sigma}$ is further defined below. The random stimulus effects are orthogonal to the random effects concerning the individuals. They are also assumed to be randomly sampled from a population distribution (of stimuli),

$$\phi_{\delta(j)} \sim \text{Normal}(0, \sigma_{\delta(j)}),$$

where 0 is the mean and $\sigma_{\delta(j)}$ is the standard deviation.

We need to specify priors for all fixed and random effects parameters as well as for the correlations among the random effects parameters.

Prior specification Study 1. We use a standard normal prior for the originality and utility effects on the drift rate,

$$\theta_{OR}, \theta_{UT} \sim \text{Normal}(0, 1),$$

.

For the remaining fixed effects we use the following weakly informative priors:

$$\theta_\delta \sim \text{Normal}(0, 1)$$

$$\mu_\beta \sim \text{Beta}(1.3, 1.3)$$

$$\mu_\alpha \sim \text{Normal}^+(0, 2)$$

$$\tau \sim \text{Uniform}(0, 300).$$

These prior distributions restrict the parameters to a plausible range. The range of the a-priori bias parameter is from 0 to 1 and the boundary separation is restricted to be positive. The non-decision time generally needs to be smaller than the RTs. We therefore use 300ms, the minimally required response time (see exclusion criterion II.), as upper bound for the prior on τ .

For all variability parameters we use the following prior,

$$\sigma_{\delta_i}, \sigma_{\delta_j}, \sigma_{OR}, \sigma_{UT}, \sigma_\alpha \sim \text{Normal}^+(0, 0.3).$$

Lastly, we place a prior on the random effects correlations concerning the individuals. The variance-covariance matrix Σ needs to be decomposed such that we can specify a prior for the correlations only. We refer to the matrix containing the random effects correlations as \mathbf{P} . Specifically, Σ can be rewritten as $\Phi \mathbf{P} \Phi$, whereby Φ is a 4x4 matrix with only the variability parameters on the diagonal, $\Phi = \text{diag}(\sigma_{\delta_i}, \sigma_{OR}, \sigma_{UT}, \sigma_\alpha)$, and \mathbf{P} is a 4x4 correlation matrix. Exemplarily, the correlation between the random originality and utility effects is expressed as $\rho_{\sigma_{OR}\sigma_{UT}}$. We place a Lewandowski-Kurowicka-Joe (LKJ) prior with the shape 3 on \mathbf{P} (Lewandowski, Kurowicka, & Joe, 2009),

$$\mathbf{P} \sim \text{LKJ}(3).$$

This prior restricts the correlations to the range -1 to 1, makes it a proper correlation matrix, and places most prior mass around 0.

Prior specification Study 2. We use informative, truncated prior distributions for the parameters where we expect a positive effect (i.e., the originality and utility effects on the drift rate),

$$\mu_{\theta_{OR}}, \mu_{\theta_{UT}} \sim \text{Normal}^+(0, 0.2),$$

where 0 is the mean and 0.2 the standard deviation. This truncated prior distribution is informed by previous research that, overall, people take into account both originality and utility when they evaluate creative ideas. It is also informed by data. In Study 1's dataset, the presence of effects of both stimulus originality and utility on the drift rate are detectable using this prior.

For the remaining fixed effects we use the following weakly informative priors:

$$\mu_{\theta_{\delta}} \sim \text{Normal}(0, 1)$$

$$\mu_{\beta} \sim \text{Beta}(1, 1)$$

$$\mu_{\alpha} \sim \text{Normal}^+(0, 2)$$

$$\tau \sim \text{Uniform}(0, 300).$$

These prior distributions again restrict the parameters to a plausible range.

For all variability parameters we use the following prior,

$$\sigma_{\delta_i}, \sigma_{\delta_j}, \sigma_{OR}, \sigma_{UT}, \sigma_{\beta}, \sigma_{\alpha} \sim \text{Normal}^+(0, 0.3).$$

This prior is again informed by previous analyses on Study 1's dataset. Lastly, we need to place a prior on the random effects correlations concerning the individuals. Here, $\mathbf{\Sigma}$ can again be rewritten as $\mathbf{\Phi P \Phi}$, whereby $\mathbf{\Phi}$ is a 5x5 matrix with only the variability parameters on the diagonal, $\mathbf{\Phi} = \text{diag}(\sigma_{\delta_i}, \sigma_{OR}, \sigma_{UT}, \sigma_{\alpha}, \sigma_{\beta})$, and \mathbf{P} is a 5x5 correlation matrix. We again place a Lewandowski-Kurowicka-Joe (LKJ) prior with the shape parameter 3 on \mathbf{P} (Lewandowski et al., 2009),

$$\mathbf{P} \sim \text{LKJ}(3).$$

Appendix B

This Appendix contains the data cleaning procedure that we applied in Study 1 and Study 2.

Study 1.

CON-task. Before analyzing the data, we employed the following exclusion criteria. The full dataset comprised 18984 trials from 299 participants. First, we excluded data from all participants who gave the same response (either “creative” or “not creative”) in at least $57/64$ ($\approx 90\%$) of the trials. This step removed data from 2 individuals. We then removed the first two trials to account for the fact that participants needed time to get acquainted with the task (= 594 trials). We also excluded all trials with response times greater than 6 seconds and less than 0.3 seconds to exclude unreasonably fast and slow responses. In this step, 338 trials were excluded. Finally, we excluded data from individuals with fewer than 47 ($\approx 3/4$) remaining trials. This last step removed data from 14 participants.

AUT. $n = 55$ of the participants who completed the CON-task did not complete the AUT. We cleaned the AUT data by removing data from all participants with less than 90 percent valid responses. For example, invalid responses are those where participants responded with associations rather than uses (e.g., “rectangular” as response for the object “brick”). We additionally removed all responses that both raters had scored as invalid. As performance indicators, we used the mean originality and mean utility score across raters and responses.

Study 2.

CON-task. The full dataset comprised 11220 trials and 176 participants. First, we excluded 10 participants for giving the same response in at least $57/64$ ($\approx 90\%$) of the trials. We then removed the first two trials for each participant (= 332 trials) and all trials with

RTs greater than 6 seconds and less than 0.3 seconds (= 585 trials). Finally, data from 14 participants were excluded because they had fewer than 47 ($\approx 3/4$) remaining trials. The sample used to estimate the DDM comprised 152 participants and 9291 trials. Although we pre-registered to include only data from Dutch native speakers, we retrospectively decided not to exclude data from non-Dutch native speakers as long as they were able to read and respond fluently in Dutch to be as inclusive as possible. Thus, we included all participants who chose to do the Dutch (rather than English) version of the experiment. Participants were required to read an instruction in Dutch in order to do the experiment in Dutch. If they chose the Dutch version, we considered them sufficiently fluent to read and respond in Dutch for our study's purposes where only a few simple phrases had to be read or written in Dutch.

AUT. In Study 2, $n = 12$ participants did not complete the AUT. We again cleaned the AUT data by removing data from all participants with less than 90 percent responses that were scored as valid by all raters. Additionally, we again removed all responses that both raters had scored as invalid. As performance indicators, we again used the mean originality and mean utility score across raters and responses.

Appendix C

Table 3

Posterior mean, standard deviation, and 95% credible interval for the correlations among random effects parameters

	Study 1				Study 2			
	Mean	SD	LB	UB	Mean	SD	LB	UB
$\rho_{\sigma_{\theta_{OR}}\sigma_{\theta_{UT}}}$	-0.44	0.08	-0.6	-0.27	-0.34	0.13	-0.58	-0.09
$\rho_{\sigma_{\delta}\sigma_{\theta_{OR}}}$	-0.17	0.07	-0.3	-0.02	-0.04	0.11	-0.24	0.18
$\rho_{\sigma_{\delta}\sigma_{\alpha}}$	-0.11	0.07	-0.24	0.03	0.02	0.10	-0.16	0.21
$\rho_{\sigma_{\theta_{UT}}\sigma_{\alpha}}$	-0.05	0.08	-0.21	0.11	-0.21	0.12	-0.43	0.02
$\rho_{\sigma_{\delta}\sigma_{\theta_{UT}}}$	0.07	0.08	-0.08	0.22	-0.09	0.11	-0.30	0.12
$\rho_{\sigma_{\theta_{OR}}\sigma_{\alpha}}$	0.31	0.07	0.17	0.44	0.24	0.11	0.02	0.44
$\rho_{\sigma_{\delta}\sigma_{\beta}}$	-	-	-	-	-0.45	0.10	-0.63	-0.24
$\rho_{\sigma_{\theta_{OR}}\sigma_{\beta}}$	-	-	-	-	-0.17	0.14	-0.44	0.12
$\rho_{\sigma_{\theta_{UT}}\sigma_{\beta}}$	-	-	-	-	0.14	0.14	-0.14	0.41
$\rho_{\sigma_{\alpha}\sigma_{\beta}}$	-	-	-	-	-0.42	0.10	-0.60	-0.21

Note. SD = standard deviation; LB = lower bound; UB = upper bound.

References

- al., A. S. et mult. (2019). *DescTools: Tools for descriptive statistics*. Retrieved from <https://cran.r-project.org/package=DescTools>
- Auguie, B. (2017). *GridExtra: Miscellaneous functions for "grid" graphics*. Retrieved from <https://CRAN.R-project.org/package=gridExtra>
- Aust, F., & Barth, M. (2020). *papaja: Create APA manuscripts with R Markdown*. Retrieved from <https://github.com/crsh/papaja>
- Bache, S. M., & Wickham, H. (2020). *Magrittr: A forward-pipe operator for r*. Retrieved from <https://CRAN.R-project.org/package=magrittr>
- Barron, F. (1955). The disposition toward originality. *The Journal of Abnormal and Social Psychology*, 51(3), 478–485. <https://doi.org/10.1037/h0048073>
- Basadur, M. (1995). Optimal ideation-evaluation ratios. *Creativity Research Journal*, 8(1), 63–75. https://doi.org/10.1207/s15326934crj0801_5
- Bates, D., & Maechler, M. (2019). *Matrix: Sparse and dense matrix classes and methods*. Retrieved from <https://CRAN.R-project.org/package=Matrix>
- Benedek, M., Fink, A., & Neubauer, A. C. (2006). Enhancement of ideational fluency by means of computer-based training. *Creativity Research Journal*, 18(3), 317–328. https://doi.org/10.1207/s15326934crj1803_7
- Boden, M. A. (2004). *The creative mind: Myths and mechanisms* (2nd edition). New York, NY: Routledge.
- Bürkner, P.-C. (2017). brms: An R package for Bayesian multilevel models using Stan. *Journal of Statistical Software*, 80(1), 1–28. <https://doi.org/10.18637/jss.v080.i01>

Bürkner, P.-C. (2018). Advanced Bayesian Multilevel Modeling with the R Package brms.

The R Journal, 10(1), 395. <https://doi.org/10.32614/RJ-2018-017>

Caroff, X., & Besançon, M. (2008). Variability of creativity judgments. *Learning and*

Individual Differences, 18(4), 367–371. <https://doi.org/10.1016/j.lindif.2008.04.001>

Carpenter, B., Gelman, A., Hoffman, M. D., Lee, D., Goodrich, B., Betancourt, M., ...

Riddell, A. (2017). Stan: A Probabilistic Programming Language. *Journal of Statistical Software*, 76(1). <https://doi.org/10.18637/jss.v076.i01>

Casner-Lotto, J., & Barrington, L. (2006). *Are they really ready to work? Employers'*

perspectives on the basic knowledge and applied skills of new entrants to the 21st century U.S. Workforce. Washington, DC: Partnership for 21st Century Skills.

de Buissonjé, D. R., Ritter, S. M., de Bruin, S., ter Horst, J. M.-L., & Meeldijk, A. (2017).

Facilitating creative idea selection: The combined effects of self-affirmation, promotion focus and positive affect. *Creativity Research Journal*, 29(2), 174–181. <https://doi.org/10.1080/10400419.2017.1303308>

Dickey, J. M. (1971). The weighted likelihood ratio, linear hypotheses on normal location

parameters. *The Annals of Mathematical Statistics*, 204–223.

Diedrich, J., Benedek, M., Jauk, E., & Neubauer, A. C. (2015). Are creative ideas novel and

useful? *Psychology of Aesthetics, Creativity, and the Arts*, 9(1), 35–40.

<https://doi.org/10.1037/a0038688>

Eddelbuettel, D., & Balamuta, J. J. (2017). Extending extitR with extitC++: A Brief

Introduction to extitRcpp. *PeerJ Preprints*, 5, e3188v1.

<https://doi.org/10.7287/peerj.preprints.3188v1>

Eddelbuettel, D., & François, R. (2011). Rcpp: Seamless R and C++ integration. *Journal of*

- Statistical Software*, 40(8), 1–18. <https://doi.org/10.18637/jss.v040.i08>
- Efron, B., & Morris, C. (1977). Stein's Paradox in Statistics. *Scientific American*, 236(5), 119–127. <https://doi.org/10.1038/scientificamerican0577-119>
- Forthmann, B., Jendryczko, D., Scharfen, J., Kleinkorres, R., Benedek, M., & Holling, H. (2019). Creative ideation, broad retrieval ability, and processing speed: A confirmatory study of nested cognitive abilities. *Intelligence*, 75, 59–72. <https://doi.org/10.1016/j.intell.2019.04.006>
- Gabry, J., Simpson, D., Vehtari, A., Betancourt, M., & Gelman, A. (2019). Visualization in bayesian workflow. *J. R. Stat. Soc. A*, 182(2), 389–402. <https://doi.org/10.1111/rssa.12378>
- Gelman, A., & Rubin, D. B. (1992). Inference from Iterative Simulation Using Multiple Sequences. *Statistical Science*, 7(4), 457–472. <https://doi.org/10.1214/ss/1177011136>
- Genz, A., & Bretz, F. (2009). *Computation of multivariate normal and t probabilities*. Heidelberg: Springer-Verlag.
- Guilford, J. P. (1967). *The nature of human intelligence*. New York, NY: McGraw-Hill.
- Haaf, J. M., & Rouder, J. N. (2017). Developing constraint in bayesian mixed models. *Psychological Methods*, 22, 779–798. <https://doi.org/10.1037/met0000156>
- Haaf, J. M., & Rouder, J. N. (2018). Some do and some don't? Accounting for variability of individual difference structures. *Psychonomic Bulletin & Review*. <https://doi.org/10.3758/s13423-018-1522-x>
- Hennessey, B. A., & Amabile, T. M. (2010). Creativity. *Annual Review of Psychology*, 61(1), 569–598. <https://doi.org/10.1146/annurev.psych.093008.100416>

- Herman, A., & Reiter-Palmon, R. (2011). The effect of regulatory focus on idea generation and idea evaluation. *Psychology of Aesthetics, Creativity, and the Arts*, 5(1), 13–20.
<https://doi.org/10.1037/a0018587>
- Hunter, S., Friedrich, T., Bedell, K., & Mumford, M. (2006). Creative thought in real-world innovation. *Serbian Journal of Management*, 1(1), 29–39.
- IBM. (2010). *Capitalizing on complexity*. IBM Global CEO Study.
- Jackson, C. H. (2011). Multi-state models for panel data: The msm package for R. *Journal of Statistical Software*, 38(8), 1–29. Retrieved from <http://www.jstatsoft.org/v38/i08/>
- Klugkist, I., & Hoijtink, H. (2007). The Bayes factor for inequality and about equality constrained models. *Computational Statistics & Data Analysis*, 51(12), 6367–6379.
<https://doi.org/10.1016/j.csda.2007.01.024>
- Klugkist, I., Laudy, O., & Hoijtink, H. (2005). Inequality Constrained Analysis of Variance: A Bayesian Approach. *Psychological Methods*, 10(4), 477–493.
<https://doi.org/10.1037/1082-989X.10.4.477>
- Krypotos, A.-M., Beckers, T., Kindt, M., & Wagenmakers, E.-J. (2015). A Bayesian hierarchical diffusion model decomposition of performance in Approach–Avoidance Tasks. *Cognition and Emotion*, 29(8), 1424–1444.
<https://doi.org/10.1080/02699931.2014.985635>
- Lee, M. D. (2011). How cognitive modeling can benefit from hierarchical Bayesian models. *Journal of Mathematical Psychology*, 55(1), 1–7.
<https://doi.org/10.1016/j.jmp.2010.08.013>
- Lewandowski, D., Kurowicka, D., & Joe, H. (2009). Generating random correlation matrices based on vines and extended onion method. *Journal of Multivariate Analysis*, 100(9),

- 1989–2001. <https://doi.org/10.1016/j.jmva.2009.04.008>
- Loewenstein, J., & Mueller, J. (2016). Implicit theories of creative ideas: How culture guides creativity assessments. *Academy of Management Discoveries*, 2(4), 320–348. <https://doi.org/10.5465/amd.2014.0147>
- Matzke, D., & Wagenmakers, E.-J. (2009). Psychological interpretation of the ex-Gaussian and shifted Wald parameters: A diffusion model analysis. *Psychonomic Bulletin & Review*, 16(5), 798–817. <https://doi.org/10.3758/PBR.16.5.798>
- Meyer, D. E., & Schvaneveldt, R. W. (1971). Facilitation in recognizing pairs of words: Evidence of a dependence between retrieval operations. *Journal of Experimental Psychology*, 90(2), 227–234. <https://doi.org/10.1037/h0031564>
- Morey, R. D., & Rouder, J. N. (2018). *BayesFactor: Computation of bayes factors for common designs*. Retrieved from <https://CRAN.R-project.org/package=BayesFactor>
- Mueller, J., Melwani, S., Loewenstein, J., & Deal, J. J. (2018). Reframing the decision-makers' dilemma: Towards a social context model of creative idea recognition. *Academy of Management Journal*, 61(1), 94–110. <https://doi.org/10.5465/amj.2013.0887>
- Müller, K. (2017). *Here: A simpler way to find your files*. Retrieved from <https://CRAN.R-project.org/package=here>
- Müller, K., & Wickham, H. (2020). *Tibble: Simple data frames*. Retrieved from <https://CRAN.R-project.org/package=tibble>
- Plummer, M., Best, N., Cowles, K., & Vines, K. (2006). CODA: Convergence diagnosis and output analysis for mcmc. *R News*, 6(1), 7–11. Retrieved from <https://journal.r-project.org/archive/>

- Randolph, M. (2019). *That will never work: The birth of Netflix and the amazing life of an idea*. New York, NY: Hachette Book Group.
- Ratcliff, R. (1978). A theory of memory retrieval. *Psychological Review*, 85(2), 59–108.
<https://doi.org/10.1037/0033-295x.85.2.59>
- Ratcliff, R., & McKoon, G. (2008). The Diffusion Decision Model: Theory and Data for Two-Choice Decision Tasks. *Neural Computation*, 20(4), 873–922.
<https://doi.org/10.1162/neco.2008.12-06-420>
- R Core Team. (2019). *R: A language and environment for statistical computing*. Vienna, Austria: R Foundation for Statistical Computing. Retrieved from
<https://www.R-project.org/>
- Revelle, W. (2018). *Psych: Procedures for psychological, psychometric, and personality research*. Evanston, Illinois: Northwestern University. Retrieved from
<https://CRAN.R-project.org/package=psych>
- Rietzschel, E. F., Nijstad, B. A., & Stroebe, W. (2010). The selection of creative ideas after individual idea generation: Choosing between creativity and impact. *British Journal of Psychology*, 101(1), 47–68. <https://doi.org/10.1348/000712609X414204>
- Rouder, J. N., & Lu, J. (2005). An introduction to Bayesian hierarchical models with an application in the theory of signal detection. *Psychonomic Bulletin & Review*, 12(4), 573–604. <https://doi.org/10.3758/BF03196750>
- Runco, M. A., & Acar, S. (2012). Divergent thinking as an indicator of creative potential. *Creativity Research Journal*, 24(1), 66–75.
<https://doi.org/10.1080/10400419.2012.652929>
- Runco, M. A., & Basadur, M. (1993). Assessing ideational and evaluative skills and creative

- styles and attitudes. *Creativity and Innovation Management*, 2(3).
<https://doi.org/10.1109/iemc.1990.201291>
- Runco, M. A., & Charles, R. E. (1993). Judgments of originality and appropriateness as predictors of creativity. *Personality and Individual Differences*, 15(5), 537–546.
[https://doi.org/10.1016/0191-8869\(93\)90337-3](https://doi.org/10.1016/0191-8869(93)90337-3)
- Runco, M. A., & Jaeger, G. J. (2012). The Standard Definition of Creativity. *Creativity Research Journal*, 24(1), 92–96. <https://doi.org/10.1080/10400419.2012.650092>
- Runco, M. A., & Smith, W. R. (1992). Interpersonal and intrapersonal evaluations of creative ideas. *Personality and Individual Differences*, 13(3), 295–302.
[https://doi.org/10.1016/0191-8869\(92\)90105-x](https://doi.org/10.1016/0191-8869(92)90105-x)
- Said-Metwaly, S., Fernández-Castilla, B., Kyndt, E., & Van den Noortgate, W. (2020). Testing conditions and creative performance: Meta-analyses of the impact of time limits and instructions. *Psychology of Aesthetics, Creativity, and the Arts*, 14(1), 15–38. <https://doi.org/10.1037/aca0000244>
- Simonton, D. K. (2018). Defining Creativity: Don't We Also Need to Define What Is *Not* Creative? *The Journal of Creative Behavior*, 52(1), 80–90.
<https://doi.org/10.1002/jocb.137>
- Singmann, H. (2018, September 6). Diffusion/Wiener Model Analysis with brms – Part III: Hypothesis Tests of Parameter Estimates. Retrieved from <http://singmann.org/wiener-model-analysis-with-brms-part-iii/>
- Stan Development Team. (2019a). RStan: The R interface to Stan. Retrieved from <http://mc-stan.org/>
- Stan Development Team. (2019b). StanHeaders: Headers for the R interface to Stan.

Retrieved from <http://mc-stan.org/>

Stein, M. I. (1953). Creativity and Culture. *The Journal of Psychology*, *36*(2), 311–322.

<https://doi.org/10.1080/00223980.1953.9712897>

Sternberg, R. J., & Lubart, T. I. (1996). Investing in creativity, *51*(7), 677–688.

Stevenson, C. E., Baas, M., & van der Maas, H. L. J. (2021). Creative or not: The originality-utility tradeoff in divergent thinking.

Tuerlinckx, F. (2004). The efficient computation of the cumulative distribution and probability density functions in the diffusion model. *Behavior Research Methods, Instruments, & Computers*, *36*(4), 702–716. <https://doi.org/10.3758/BF03206552>

Vandekerckhove, J., Tuerlinckx, F., & Lee, M. D. (2011). Hierarchical diffusion models for two-choice response times. *Psychological Methods*, *16*(1), 44–62. <https://doi.org/10.1037/a0021765>

Vehtari, A., Gelman, A., Simpson, D., Carpenter, B., & Bürkner, P.-C. (2020, January 16). Rank-normalization, folding, and localization: An improved \widehat{R} for assessing convergence of MCMC. Retrieved February 4, 2020, from <http://arxiv.org/abs/1903.08008>

Wabersich, D., & Vandekerckhove, J. (2014). The RWiener package: An R package providing distribution functions for the Wiener diffusion model. *The R Journal*, *6*(1), 49. <https://doi.org/10.32614/RJ-2014-005>

Wagenmakers, E.-J. (2009). Methodological and empirical developments for the Ratcliff diffusion model of response times and accuracy. *European Journal of Cognitive Psychology*, *21*(5), 641–671. <https://doi.org/10.1080/09541440802205067>

Wagenmakers, E.-J., Ratcliff, R., Gomez, P., & McKoon, G. (2008). A diffusion model

- account of criterion shifts in the lexical decision task. *Journal of Memory and Language*, 58(1), 140–159. <https://doi.org/10.1016/j.jml.2007.04.006>
- Wei, T., & Simko, V. (2017). *R package "corrplot": Visualization of a correlation matrix*. Retrieved from <https://github.com/taiyun/corrplot>
- Wickham, H. (2016). *Ggplot2: Elegant graphics for data analysis*. Springer-Verlag New York. Retrieved from <https://ggplot2.tidyverse.org>
- Wickham, H. (2019). *Stringr: Simple, consistent wrappers for common string operations*. Retrieved from <https://CRAN.R-project.org/package=stringr>
- Wickham, H., François, R., Henry, L., & Müller, K. (2020). *Dplyr: A grammar of data manipulation*. Retrieved from <https://CRAN.R-project.org/package=dplyr>
- Wickham, H., & Henry, L. (2019). *Tidyr: Tidy messy data*. Retrieved from <https://CRAN.R-project.org/package=tidyr>
- Wilke, C. O. (2019). *Cowplot: Streamlined plot theme and plot annotations for 'ggplot2'*. Retrieved from <https://CRAN.R-project.org/package=cowplot>